

Structured Query Language



Domain-Specific Language used in programming and designed for managing data held in a Relational Database Management System (RDBMS)

Course Agenda



Duration	Day 1	Day 2
Start Time	Introduction to SQL	Exercise 1 Review
ST + 15 min		SQL Aggregate
ST + 30 min	SQL JOINS	Function
ST + 45 min		SQL Window
ST + 60 min	Lab 1 & Exercise 1	Function
ST + 75 Min		Lab 2 & Exercise 2

Course Progress



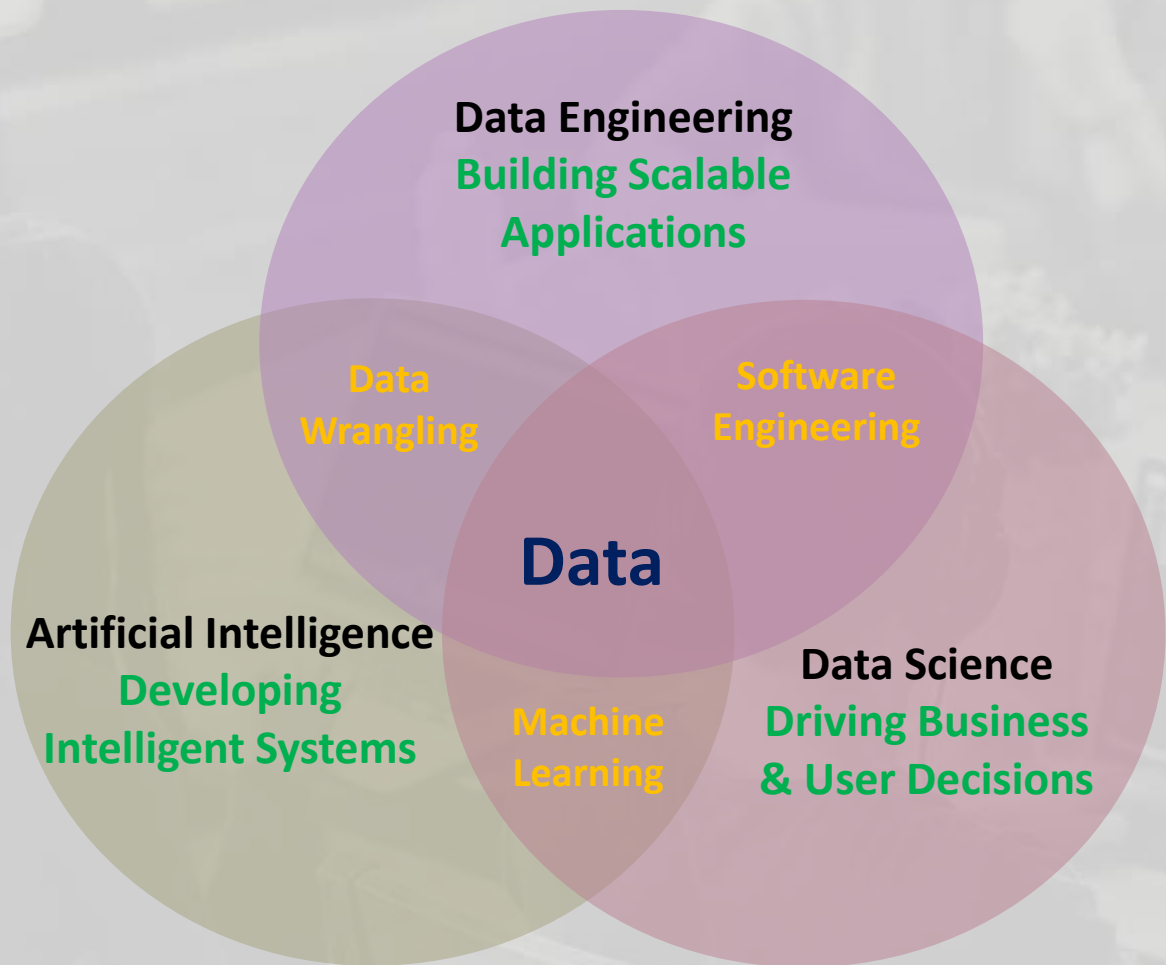
Day 1

Introduction to SQL
SQL JOINS
Lab 1 & Exercise 1

Day 2

Exercise 1 Review
SQL Aggregate Function
SQL Window Function
Lab 2 & Exercise 2

Vimalkumar Narasimman



connect me on **Linked in**



SQL



Learning Objectives

- Understand the basic concept of SQL
- Understand SELECT query execution process
- Learn SQL JOINS



Basic concepts of SQL

- ❖ **Data** are units of information
- ❖ Logically organized data in a row-and-column format is called **Records**
- ❖ Each row represents a unique record, and each column represents a field in the record
- ❖ A collection of related records are grouped and stored in a **Table**
- ❖ Tables are database objects that contain all the data in a **Database**

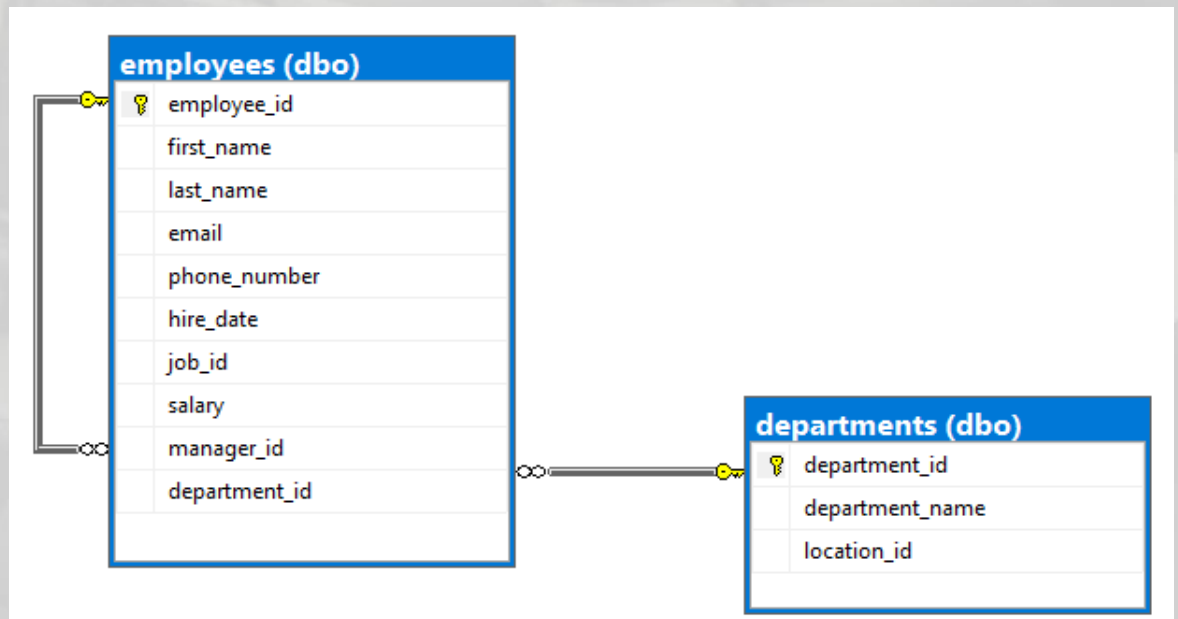
For example, a table that contains employee data for a company might contain a row for each employee and columns representing employee information such as employee number, name, address, job title, and home telephone number.



Basic concepts of SQL

- ❖ A relationship is formed by correlating rows belonging to different tables
- ❖ A table relationship is established when a child table defines a Foreign Key column that references the Primary Key column of its parent table.
- ❖ **one-to-many** is the most common relationship, and it associates a row from a parent table to multiple rows in a child table.

For example, the “Department” and “Employee” tables have a one-to-many relationship. That is, each department have many employees. But each employee comes from only one department.






Basic concepts of SQL

- ❖ **one-to-one** requires the child table Primary Key to be associated via a Foreign Key with the parent table Primary Key column.

For example, the “Employee” and “Payroll” tables have a one-to-one relationship. That is, each employee will have only one payroll.

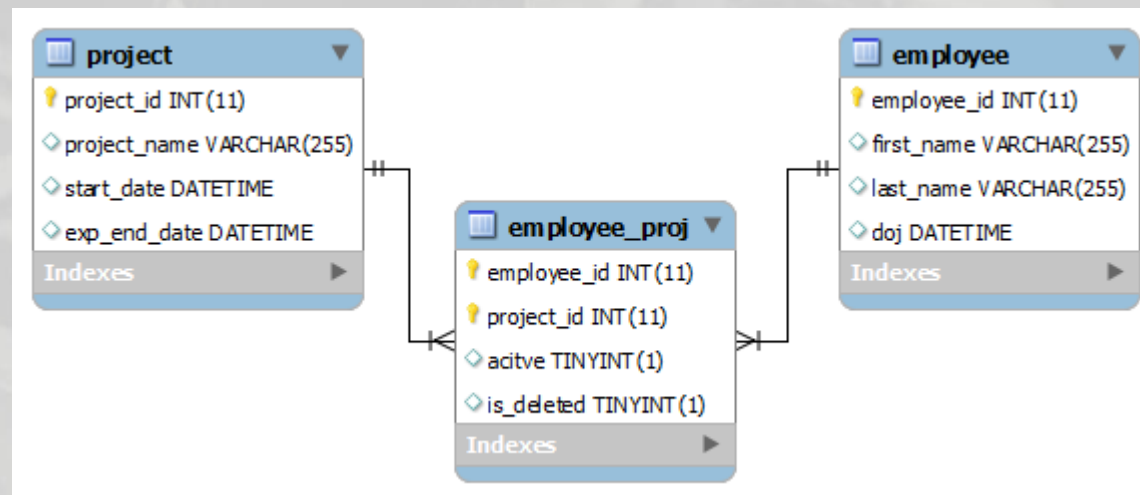


Basic concepts of SQL



- ❖ **many-to-many** requires a link table containing two Foreign Key columns that reference the two different parent tables.

For example, the “Employee” and “Project” tables have a many-to-many relationship. That is, each employee will assigned to many projects. Also each project will mapped with many employees.



Basic concepts of SQL



❖ Transitive Dependencies

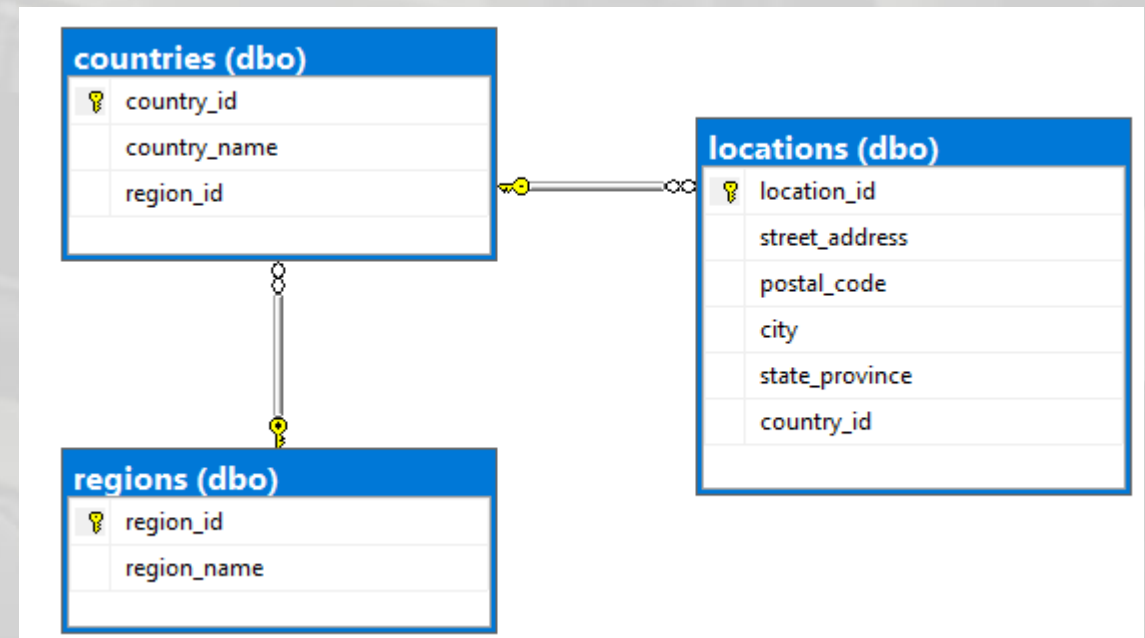
A transitive dependency exists when you have the following functional dependency pattern:

$A \rightarrow B$ and $B \rightarrow C$; therefore $A \rightarrow C$

For example, the “Regions” and “Countries” tables have relationship. the “Countries” and “Locations” tables have relationship.

Therefore the “Regions” and “Locations” tables have functional dependence

That is, each region will have location using country records





Simple SELECT statement

The SQL **SELECT** statement is used to retrieve records from one or more tables in your SQL database. The records retrieved are known as a result set.

The syntax for the SELECT statement in SQL is:

```
SELECT expressions  
FROM tables  
[WHERE conditions]  
[ORDER BY expression [ ASC | DESC ]];
```

- **Expressions:** The columns or calculations that you wish to retrieve. Use * if you wish to select all columns
- **Tables:** The tables that you wish to retrieve records from. There must be at least one table listed in the FROM clause
- **WHERE conditions:** Optional. The conditions that must be met for the records to be selected
- **ORDER BY expression:** Optional. The expression used to sort the records in the result set
- **ASC:** Optional. ASC sorts the result set in ascending order by expression. This is the default behavior, if no modifier is provided
- **DESC:** Optional. DESC sorts the result set in descending order by expression

Simple SELECT Query



supplier_id	supplier_name	city	state
100	Microsoft	Redmond	Washington
200	Google	Mountain View	California
300	Oracle	Redwood City	California
400	Kimberly-Clark	Irving	Texas
500	Tyson Foods	Springdale	Arkansas
600	SC Johnson	Racine	Wisconsin
700	Dole Food Company	Westlake Village	California
800	Flowers Foods	Thomasville	Georgia
900	Electronic Arts	Redwood City	California

```
SELECT supplier_name, city  
FROM suppliers  
WHERE supplier_id > 500  
ORDER BY supplier_name ASC, city DESC;
```

supplier_name	city
Dole Food Company	Westlake Village
Electronic Arts	Redwood City
Flowers Foods	Thomasville
SC Johnson	Racine

This example would return only the **supplier_name** and **city** fields from the **suppliers** table where the **supplier_id** value is greater than **500**. The results are sorted by **supplier_name** in ascending order and then **city** in descending order.



Simple SELECT statement Process Steps

1. Getting Data (**FROM** clause)
 2. Row Filter (**WHERE** clause)
 3. Return Expressions (**SELECT** clause)
 4. Order (**ORDER BY** clause)
- The first step in the process is the execution of the statements in FROM clause
 - After getting qualified rows, it is passed on to the WHERE clause. This evaluates every row using conditional expressions. When rows do not evaluate to true, they will be removed from the set
 - In the next step, the processor evaluates what will be printed as a result of the query, and if there are some functions to run on data like *Distinct*, *Max*, *Sqrt*, *Date*, *Lower*, etc.
 - The final processing steps of the query deal with presentation ordering the result set

JOINS



- ❖ A **SQL Join** statement is used to combine data or rows from two or more tables based on a common field between them.
- ❖ To understand this easily, let's look at the following employees and departments tables. Here, the dept_id column of the employees table is the foreign key to the departments table. Therefore, these two tables can be joined to get the combined data.

- ❖ Different types of Joins

- **INNER JOIN**
- **LEFT JOIN**
- **RIGHT JOIN**
- **FULL JOIN**

emp_id	emp_name	hire_date	dept_id	dept_id	dept_name
1	Ethan Hunt	2001-05-01	4	1	Administration
2	Tony Montana	2002-07-15	1	2	Customer Service
3	Sarah Connor	2005-10-18	5	3	Finance
4	Rick Deckard	2007-01-03	3	4	Human Resources
5	Martin Blank	2008-06-24	NULL	5	Sales


Table: **employees**Table: **departments**

JOINS




- **INNER JOIN:** Returns records that have matching values in both tables

```
SELECT emp.emp_id, emp.emp_name, emp.hire_date, dept.dept_name
FROM employees AS emp
INNER JOIN departments AS dept ON emp.dept_id = dept.dept_id
```




- **LEFT JOIN:** Returns all records from the left table, and the matched records from the right table

```
SELECT emp.emp_id, emp.emp_name, emp.hire_date, dept.dept_name
FROM employees AS emp
LEFT JOIN departments AS dept ON emp.dept_id = dept.dept_id
```




- **RIGHT JOIN:** Returns all records from the right table, and the matched records from the left table

```
SELECT emp.emp_id, emp.emp_name, emp.hire_date, dept.dept_name
FROM employees AS emp
RIGHT JOIN departments AS dept ON emp.dept_id = dept.dept_id
```



- **FULL JOIN:** Returns all records when there is a match in either left or right table

```
SELECT emp.emp_id, emp.emp_name, emp.hire_date, dept.dept_name
FROM employees AS emp
FULL OUTER JOIN departments AS dept ON emp.dept_id = dept.dept_id
```



Exercise 1 - A



The Table creation queries are present in this [link](#)

This query should return only the *employee_name* , *email*, *phone_number*, *hire_date* and *salary* fields from the *employees* table where the *hire_date* value is greater than “1990-01-01” and *salary* in between 6,000.00 and 9999.00. The results are sorted by *hire_date* in descending order and then *employee_name* in ascending order.

Note: Concatenate with *first_name* and *last_name* to get *employee_name*



Exercise 1 - B

The Table creation queries are present in this [link](#)

This query should return employee details along with dependents details with *employee_name* , *email*, *phone_number*, *hire_date*, *salary*, *dependent_name* and *relationship* fields from the *employees* and *dependents* table where the *department_name* is “Sales” or “Finance”. The results are sorted by *employee_name* and then *dependent_name* .

Note: Concatenate with *first_name* and *last_name* to get *dependent_name*

Course Progress



Day 1

Introduction to SQL
SQL JOINS
Lab 1 & Exercise 1

Day 2

Exercise 1 review
SQL Aggregate Function
SQL Window Function
Lab 2 & Exercise 2

Recap of Day 1




Data Relationships:


- One-to-One
- One-to-Many
- Many-to-Many

```
SELECT expressions  
FROM tables  
[WHERE conditions]  
[ORDER BY expression [ ASC | DESC ]];
```


INNER JOIN




LEFT JOIN



RIGHT JOIN



FULL OUTER JOIN



Day 2 Learning Objectives



- Understand the Aggregate functions
- Understand the Window Functions



Aggregate Functions

- ❖ In database management, an **aggregate function** or **aggregation function** is a function where the values of multiple rows are grouped together to form a single summary statistics

Common aggregate functions include:

- Average (i.e., arithmetic mean)
- Count
- Maximum
- Minimum
- Sum

Others Include:

- Median
- Mode
- Range
- Stddev



Simple Aggregation Query

supplier_id	supplier_name	city	state
100	Microsoft	Redmond	Washington
200	Google	Mountain View	California
300	Oracle	Redwood City	California
400	Kimberly-Clark	Irving	Texas
500	Tyson Foods	Springdale	Arkansas
600	SC Johnson	Racine	Wisconsin
700	Dole Food Company	Westlake Village	California
800	Flowers Foods	Thomasville	Georgia
900	Electronic Arts	Redwood City	California

```
SELECT  
COUNT(supplier_name) AS number_of_suppliers  
FROM suppliers
```

number_of_suppliers
9

This example would return COUNT of the **supplier_name**

GROUP BY and HAVING statement



The GROUP BY clause instructs the RDBMS to group the data and then perform the aggregate (function) on each group rather than on the entire result set.

Aside from the aggregate calculation statements, every column in your SELECT statement must be present in the GROUP BY clause.

The HAVING clause is used in combination with the GROUP BY clause to restrict the groups of returned rows to only those condition is TRUE

The GROUP BY clause must come after any WHERE clause and before any ORDER BY clause.

The syntax for the GROUP BY and HAVING statement in SQL is:

```
SELECT expressions
FROM tables
[WHERE conditions]
[GROUP BY expressions
    [HAVING conditions]]
[ORDER BY expression [ ASC | DESC ]];
```

Simple Aggregation Query with GROUP BY




supplier_id	supplier_name	city	state
100	Microsoft	Redmond	Washington
200	Google	Mountain View	California
300	Oracle	Redwood City	California
400	Kimberly-Clark	Irving	Texas
500	Tyson Foods	Springdale	Arkansas
600	SC Johnson	Racine	Wisconsin
700	Dole Food Company	Westlake Village	California
800	Flowers Foods	Thomasville	Georgia
900	Electronic Arts	Redwood City	California

```
SELECT state,  
COUNT(supplier_name) AS number_of_suppliers  
FROM suppliers  
GROUP BY state  
ORDER BY state
```



state	
Arkansas	
California	
Georgia	
Texas	
Washington	
Wisconsin	



state	number_of_suppliers
Arkansas	1
California	4
Georgia	1
Texas	1
Washington	1
Wisconsin	1

Simple Aggregation Query with GROUP BY and HAVING



supplier_id	supplier_name	city	state
100	Microsoft	Redmond	Washington
200	Google	Mountain View	California
300	Oracle	Redwood City	California
400	Kimberly-Clark	Irving	Texas
500	Tyson Foods	Springdale	Arkansas
600	SC Johnson	Racine	Wisconsin
700	Dole Food Company	Westlake Village	California
800	Flowers Foods	Thomasville	Georgia
900	Electronic Arts	Redwood City	California

```
SELECT state,  
COUNT(supplier_name) AS number_of_suppliers  
FROM suppliers  
GROUP BY state HAVING COUNT(supplier_name) > 1  
ORDER BY state
```

state	
Arkansas	
California	
Georgia	
Texas	
Washington	
Wisconsin	

state	number_of_suppliers
Arkansas	1
California	4
Georgia	1
Texas	1
Washington	1
Wisconsin	1

state	number_of_suppliers
California	4



SELECT statement Process Steps

1. Getting Data (**FROM, JOIN** clause)
2. Row Filter (**WHERE** clause)
3. Grouping (**GROUP BY** clause)
4. Group Filter (**HAVING** clause)
5. Return Expressions (**SELECT** clause)
6. Order & Paging (**ORDER BY** clause & **LIMIT / OFFSET**)



Exercise 2 - A

The Table creation queries are present in this [link](#)

This query should return all available *region_name*, *country_name*, *city*, *department_name* and *number_of_employee* fields from the relevant tables. The results are sorted by *country_name* and then *dependent_name*.

Note: *number_of_employee* can be zero if no employee is working in particular *region_name*, *country_name*, *city* **and** *dependent_name*

Exercise 2 - B



The Table creation queries are present in this [link](#)

This query should return all available *region_name*, *country_name*, *department_name*, *average_employee_salary* and *recent_hiring_date* fields from the relevant tables. The results are sorted by *recent_hiring_date*.



Window Functions

- ❖ In SQL, a **window function** or **analytic function** is a function which uses values from one or multiple rows to return a value for each row
- ❖ This contrasts with an aggregate function, which returns a single value for multiple rows
- ❖ the OVER clause defines a window or user-specified set of rows within a query result set
- ❖ A window function then computes a value for each row in the window
- ❖ You can use the OVER clause with functions to compute aggregated values such as moving averages, cumulative aggregates, running totals, or a top N per group results



Types of Window Functions

➤ Aggregate Window Functions

`SUM()`, `MAX()`, `MIN()`, `AVG()`, `COUNT()`

➤ Ranking Window Functions

`RANK()`, `DENSE_RANK()`, `ROW_NUMBER()`, `NTILE()`

➤ Analytic Window Functions

`LAG()`, `LEAD()`, `FIRST_VALUE()`, `LAST_VALUE()`

```
WINDOW_FUNCTION(expression)
OVER ( [ <PARTITION BY clause> ] [ <ORDER BY clause> ] [ <ROW or RANGE clause> ] )
```

- ✓ **PARTITION BY** that divides the query result set into partitions.
- ✓ **ORDER BY** that defines the logical order of the rows within each partition of the result set.
- ✓ **ROWS/RANGE** that limits the rows within the partition by specifying start and end points within the partition.

Aggregate Window Query



```
SELECT order_id, order_date, customer_name, city, order_amount  
      ,SUM(order_amount) OVER(PARTITION BY city) as grand_total  
FROM [dbo].[Orders]
```

	order_id	order_date	customer_name	city	order_amount	grand_total
1	1002	2017-04-02	David Jones	Arlington	20000.00	37000.00
2	1007	2017-04-10	Andrew Smith	Arlington	15000.00	37000.00
3	1008	2017-04-11	David Brown	Arlington	2000.00	37000.00
4	1001	2017-04-01	David Smith	GuildFord	10000.00	50500.00
5	1006	2017-04-06	Paum Smith	GuildFord	25000.00	50500.00
6	1004	2017-04-04	Michael Smith	GuildFord	15000.00	50500.00
7	1010	2017-04-25	Peter Smith	GuildFord	500.00	50500.00
8	1005	2017-04-05	David Williams	Shalford	7000.00	13000.00
9	1003	2017-04-03	John Smith	Shalford	5000.00	13000.00
10	1009	2017-04-20	Robert Smith	Shalford	1000.00	13000.00

Ranking Window Query



```
SELECT order_id,order_date,customer_name,city, order_amount,  
ROW_NUMBER() OVER(PARTITION BY city ORDER BY order_amount DESC) [row_number]  
FROM [dbo].[Orders]
```

	order_id	order_date	customer_name	city	order_amount	row_number
1	1002	2017-04-02	David Jones	Arlington	20000.00	1
2	1007	2017-04-10	Andrew Smith	Arlington	15000.00	2
3	1008	2017-04-11	David Brown	Arlington	2000.00	3
4	1006	2017-04-06	Paum Smith	GuildFord	25000.00	1
5	1004	2017-04-04	Michael Smith	GuildFord	15000.00	2
6	1001	2017-04-01	David Smith	GuildFord	10000.00	3
7	1010	2017-04-25	Peter Smith	GuildFord	500.00	4
8	1005	2017-04-05	David Williams	Shalford	7000.00	1
9	1003	2017-04-03	John Smith	Shalford	5000.00	2
10	1009	2017-04-20	Robert Smith	Shalford	1000.00	3

Analytic Window Query



```
SELECT order_id,order_date,customer_name,city, order_amount,  
FIRST_VALUE(order_date) OVER(PARTITION BY city ORDER BY city) first_order_date,  
LAST_VALUE(order_date) OVER(PARTITION BY city ORDER BY city) last_order_date  
FROM [dbo].[Orders]
```

	order_id	order_date	customer_name	city	order_amount	first_order_date	last_order_date
1	1002	2017-04-02	David Jones	Arlington	20000.00	2017-04-02	2017-04-11
2	1007	2017-04-10	Andrew Smith	Arlington	15000.00	2017-04-02	2017-04-11
3	1008	2017-04-11	David Brown	Arlington	2000.00	2017-04-02	2017-04-11
4	1001	2017-04-01	David Smith	GuildFord	10000.00	2017-04-01	2017-04-25
5	1006	2017-04-06	Paum Smith	GuildFord	25000.00	2017-04-01	2017-04-25
6	1004	2017-04-04	Michael Smith	GuildFord	15000.00	2017-04-01	2017-04-25
7	1010	2017-04-25	Peter Smith	GuildFord	500.00	2017-04-01	2017-04-25
8	1005	2017-04-05	David Williams	Shalford	7000.00	2017-04-05	2017-04-20
9	1003	2017-04-03	John Smith	Shalford	5000.00	2017-04-05	2017-04-20
10	1009	2017-04-20	Robert Smith	Shalford	1000.00	2017-04-05	2017-04-20



Exercise 2 - C

The Table creation queries are present in this [link](#)

This query should return all available *region_name*, *country_name*, *department_name*, *number_of_employee*, *first_joined_employee_name* and *last_joined_employee_name* fields from the relevant tables. The results are sorted by *country_name* and then *dependent_name*.

Note: *number_of_employee* can be zero, *first_joined_employee_name* and *last_joined_employee_name* can be null if no employee is working in particular *region_name*, *country_name* **and** *dependent_name*

Exercise 2 - D



The Table creation queries are present in this [link](#)

This query should return all available *region_name*, *country_name*, *department_name*, *average_employee_salary* and *total_employee_salary* by each *region*, *country* and *department* fields from the relevant tables. The results are sorted by *recent_hiring_date*.



Reference Link

- [SQL Tutorial](#)
- [SELECT statement Process Step](#)
- [Aggregate Functions](#)
- [Window Functions](#)

THANK YOU!